

# On the Time Complexity of Common All-Reduce Methods

Kun Yuan

(Internal Reference Only)

September 7, 2019

# Outline

1. Training framework for distributed deep learning
2. Time complexity for each all-reduce method
3. Decentralized algorithms

# Outline

1. Training framework for distributed deep learning
2. Time complexity for each all-reduce method
3. Decentralized algorithms

# Empirical Risk Minimization

- The empirical risk minimization (ERM) problem is

$$\min_w F(w) = \frac{1}{N} \sum_{n=1}^N Q(w; x_n)$$

where  $\{x_n\}_{n=1}^N$  is the dataset;  $Q(w; x_n)$  is some cost function

- The above ERM problem can be rewritten as

$$\min_w F(w) = \frac{1}{P} \sum_{p=1}^P F_p(w), \quad \text{where} \quad F_p(w) = \frac{1}{L} \sum_{n=1}^L Q(w; x_{p,n}) \quad (1)$$

where  $\{x_{p,n}\}_{n=1}^L$  is the data assigned to node  $p$ , and  $N = LP$ .

# Gradient Descent Method

- The gradient descent method to solve problem (1) is

$$w^{k+1} = w^k - \frac{\alpha}{P} \sum_{p=1}^P \nabla F_p(w^k),$$

$$\text{where } \nabla F_p(w^k) = \frac{1}{L} \sum_{n=1}^L \nabla Q(w; x_{p,n})$$

where  $\alpha$  is the step-size.

- The computation of the true gradient  $\nabla F_p(w^k)$  is time-consuming.

# Stochastic Gradient Descent

- We compute an approximate gradient for  $\nabla F_p(w)$ , i.e.,  $\widehat{\nabla F_p}(w)$ :

$$w^{k+1} = w^k - \frac{\alpha}{P} \sum_{p=1}^P \widehat{\nabla F_p}(w^k),$$

$$\text{where } \widehat{\nabla F_p}(w^k) = \frac{1}{B} \sum_{n \in \mathcal{B}} \nabla Q(w; x_{p,n})$$

where  $\mathcal{B}$  is a small batch of index, and  $B = |\mathcal{B}|$ .

- The batch is usually sampled without replacement.
- SGD is a standard training algorithm for deep learning.
- More advanced training methods: ADAM and its variant

# Training framework for deep learning: Parameter Server

---

## Algorithm 1 Distributed training framework: Parameter Server

---

**Setting:**  $P$  nodes; Partitioned data set  $\{\{x_{1,n}\}_{n=1}^{N/K}, \dots, \{x_{K,n}\}_{n=1}^{N/K}\}$

**For**  $k = 1, 2, 3 \dots$  **until meet a criterion:**

**All**  $(P - 1)$  **slave nodes do in parallel:**

        Pull  $w^k$  from the master node  $P$ ;

        Sample a batch of local data  $\{x_{p,\mathcal{B}}\}$ ;

        Evaluate a local stochastic gradient  $\widehat{\nabla F}_p(w^k; \{x_{p,\mathcal{B}}\})$ ;

        Push  $\widehat{\nabla F}_p(w^k; \{x_{p,\mathcal{B}}\})$  to the master node  $P$ ;

**The master node  $P$  does:**

        Reduce all received gradients:  $\widehat{\nabla F}(w^k) = \frac{1}{P} \sum_{n=1}^P \widehat{\nabla F}_p(w^k; \{x_{p,\mathcal{B}}\})$

        Update the variable:  $w^{k+1} = w^k - \alpha \widehat{\nabla F}(w^k)$

**Output:**  $w^k$

---

# Training framework for deep learning: Parameter Server

- The all-reduce procedure in PS bases on “reduce  $\rightarrow$  broadcast”
- Has to distinguish master node from slave nodes
- Four implementations on “reduce  $\rightarrow$  broadcast”
  - Star implementation
  - Tree implementation
  - Cycle implementation
  - BytePS implementation
- Will discuss each implementation later



# Training framework for deep learning: Ring-Allreduce

---

## Algorithm 1 Distributed training framework: Ring-Allreduce

---

**Setting:**  $P$  nodes; Partitioned data set  $\{\{x_{1,n}\}_{n=1}^{N/K}, \dots, \{x_{K,n}\}_{n=1}^{N/K}\}$

**For**  $k = 1, 2, 3 \dots$  **until meet a criterion:**

**Each node does in parallel:**

        Sample a batch of local data  $\{x_{p,\mathcal{B}}\}$ ;

        Evaluate a local stochastic gradient  $\widehat{\nabla F}_p(w^k; \{x_{p,\mathcal{B}}\})$ ;

        Run ring-allreduce and get  $\widehat{\nabla F}(w^k) = \frac{1}{P} \sum_{n=1}^P \widehat{\nabla F}_p(w^k; \{x_{p,\mathcal{B}}\})$ ;

        Update the variable:  $w^{k+1} = w^k - \alpha \widehat{\nabla F}(w^k)$

**Output:**  $w^k$

---

Important note: after ring-allreduce, every node will get  $\widehat{\nabla F}(w^k)$

# Training framework for deep learning: Ring-Allreduce

- Ring-allreduce framework is fundamentally different from PS
- In ring-allreduce:
  - Every node will get  $\widehat{\nabla F}(w^k)$  after ring-allreduce
  - Every node will perform the SGD update
- In parameter server:
  - The master node will get  $\widehat{\nabla F}(w^k)$  after reduce
  - The master node will perform the SGD update

# Outline

1. Training framework for distributed deep learning
2. Time complexity for each all-reduce method
3. Decentralized algorithms

## Star All-reduce




Figure: Illustration of the star all-reduce.

## Star All-reduce

- We assume some parameters
  - $M$ : the dimension of the vector
  - $B$ : the bytes used to store each element in the vector
  - $t_m$ : the time used to transmit a unit byte
  - $t_c$ : the time used to establish one connection
- The time complexity for one star-allreduce step:

$$T_{\text{star}} = 2(P - 1)(MBt_m + t_c)$$

- Can be implemented in an asynchronous manner; has theoretical guarantee

## Tree All-reduce




Figure: Illustration of the tree all-reduce.


## Tree All-reduce

- The time complexity for one tree-allreduce step:
  - Time spent in each level of the tree:  $2(MBt_m + t_c)$
  - There are  $\log(P)$  levels in total
  - There are a “push” step and a “pull” step

$$T_{\text{tree}} = 4 \log(P)(MBt_m + t_c)$$

- Hard to be extended to the asynchronous version

## Cycle All-reduce



**Figure:** Illustration of the cycle all-reduce. In the figure, all information are merged to node 1.



## Cycle All-reduce

- The time complexity for one cycle-allreduce step:
  - Time spent in each step:  $MBt_m + t_c$
  - There are  $\log(P)$  steps in total
  - There are a “reduce” step and a “broadcast” step

$$T_{\text{cycle}} = 2 \log(P)(MBt_m + t_c)$$

- Hard to be extended to the asynchronous version

# BytePS All-reduce




Figure: Illustration of the BytePS all-reduce.

# BytePS All-reduce

- The time complexity for one BytePS-allreduce step:

$$T_{\text{BytePS}} = 2(MBt_m + P^2t_c)$$

- Can be improved to better complexity (but with more complicated ops.)

$$T_{\text{BytePS}} = 2(MBt_m + Pt_c)$$

- Can be extended to the asynchronous version (not quite sure)
- Require more (but cheap) servers

# Ring-allreduce

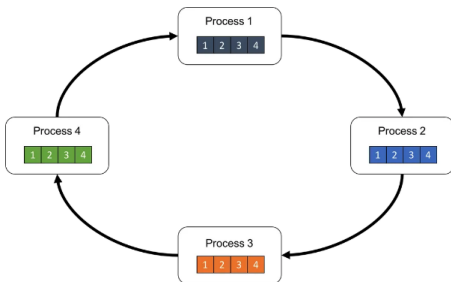


Fig.2 Example of a process ring

**Figure:** Illustration of the ring-allreduce.

# Ring-allreduce

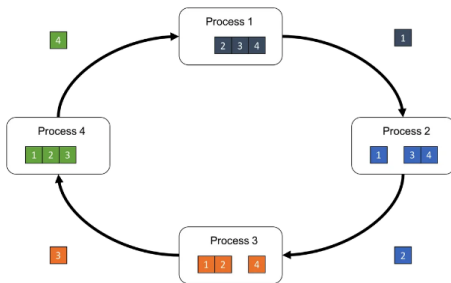


Figure: Illustration of the ring-allreduce.

# Ring-allreduce




Figure: Illustration of the ring-allreduce.

# Ring-allreduce




Fig.5 After P-1 steps, each process has a reduced subarray.

**Figure:** Illustration of the ring-allreduce.

# Ring-allreduce

- The time complexity for one Ring-allreduce step:
  - Time spent at each step:  $\frac{MB}{P}t_m + t_c$
  - There are  $2(P - 1)$  steps in total

$$T_{\text{ring}} = 2(P - 1)\left(\frac{MB}{P}t_m + t_c\right) \approx 2(MBt_m + Pt_c)$$

- Ring-allreduce has more elegant operations compared to BytePS
- Hard to be extended to the asynchronous version



# Outline

1. Training framework for distributed deep learning
2. Time complexity for each all-reduce method
3. Decentralized algorithms

# Core Idea

- All the above all-reduce methods traverse through all nodes
- A brand new idea:

Use the local average to approximate global average asymptotically

- Further lower down the above-mentioned time complexity

## Average Consensus

- Calculate the average  $\bar{x} = \frac{1}{P} \sum_{p=1}^N x_p$  over an arbitrary graph
- Core idea: each agent does local average

$$x_p^{k+1} = \sum_{q \in \mathcal{N}_p} a_{qp} x_q^k, \quad \text{where} \quad \sum_{q \in \mathcal{N}_p} a_{qp} = 1$$

- Convergence guarantee:  $\|x_p^k - \bar{x}\| = \rho^k$  where  $\rho < 1$  (linear convergence)
- This implies  $x_p^k \rightarrow \bar{x}$  asymptotically for any  $p = 1, 2, \dots, P$ .
- Proof:

$$\begin{aligned} \|\mathbf{x}^k - \bar{\mathbf{x}}^k\| &= \|A\mathbf{x}^{k-1} - \frac{1}{n} \mathbf{1}\mathbf{1}^T \bar{\mathbf{x}}^{k-1}\| \leq \|A - \frac{1}{n} \mathbf{1}\mathbf{1}^T\| \|\mathbf{x}^{k-1} - \bar{\mathbf{x}}^{k-1}\| \\ &\leq \rho \|\mathbf{x}^{k-1} - \bar{\mathbf{x}}^{k-1}\| \end{aligned}$$

where  $\bar{\mathbf{x}}^k = \bar{\mathbf{x}}^{k-1} = \dots = \bar{\mathbf{x}}^0$ .

# Average Consensus

$$x_P^{k+1} = \sum_{q \in \mathcal{N}_P} a_{qP} x_P^k$$




Figure: Illustration of an arbitrary graph and the consensus update.

# Decentralized Stochastic Gradient Descent (Diffusion)

- Recall the standard SGD:

$$w^{k+1} = w^k - \frac{\alpha}{P} \sum_{p=1}^P \widehat{\nabla F_p}(w^k),$$

$$\text{where } \widehat{\nabla F_p}(w^k) = \frac{1}{B} \sum_{n \in \mathcal{B}} \nabla Q(w^k; x_{p,n})$$

- Diffusion:

$$w_p^{k+1} = \sum_{q \in \mathcal{N}_p} a_{qp} \left( w_p^k - \alpha \widehat{\nabla F_p}(w^k) \right) \quad (\text{Local average})$$

# Theoretical Guarantee

- When each local cost function  $F_p(w)$  is convex

$$\limsup_{k \rightarrow \infty} \frac{1}{P} \sum_{p=1}^P \mathbb{E} \|w_p^k - w^*\|^2 = O\left(\frac{\alpha \sigma^2}{P\nu}\right), \quad (\text{SGD})$$

$$\limsup_{k \rightarrow \infty} \frac{1}{P} \sum_{p=1}^P \mathbb{E} \|w_p^k - w^*\|^2 = O\left(\frac{\alpha \sigma^2}{P\nu} + \frac{\lambda^2 \alpha^2 \sigma^2}{1 - \lambda}\right), \quad (\text{Diffusion})$$

where  $\sigma^2$  is the gradient noise,  $\lambda$  indicates the connectivity of the network

- When  $\alpha$  is sufficiently small or the network is well-connected ( $\lambda \rightarrow 0$ ),

$$\text{SGD} \approx \text{Diffusion}$$

- We can use local average per iteration without effecting the convergence!!

## Local Averaging Time Complexity

- We use the Power-2 graph.
- Star local-averaging:  $T = (MBt_m + t_s) \log P$ ; support asynchrony.
- Ring local-averaging:  $T = 2(MBt_m + \sqrt{P}t_s)$ ; not support asynchrony.
- Better than their counterparts.

## References: